




Simulation of Izhikevich's Large-Scale Neuronal Model on GPU

Nima Mohammadi
Under Supervision of Prof. Ganjtabesh



Large-Scale Modeling

- **Principle of hierarchical reductionism:** study smaller and smaller parts and attempt to understand the whole .
- While it remains uncertain whether a brain system can be understood as the interaction between independently describable subsystems, the brain does display **a hierarchy of spatial scales** with repeated structure such as molecules, synapses, neurons, microcircuits, networks, regions and systems.
- How the advent of large-scale modeling in computational neuroscience opens the possibility to study the dynamics of models **simultaneously incorporating various levels** from molecules to regions?

Large-Scale Modeling

Mathematical Modeling

- Experience shows that a model should be **as simple as possible** in order to be **tractable** (possible to analyze, easy to do computations with) and in order to make strong statements about the physical system being modeled.

"It can scarcely be denied that the supreme goal of all theory is to make the irreducible basic elements as simple and as few as possible without having to surrender the adequate representation of a single datum of experience."

-Einstein

- With **increased complexity** of model, **uncertainty** of modeling results increases. Also the **explanatory power** of model is lost. Of course the degree of simplicity that is achievable is dependent on the scientific question posed.

- It seems that a higher level model should be composed of component models at a lower level, and in turn must be much more complex and have more parameters than a model of a neuron. If true, one might ask **what is the proper lowest level?** After all, the deeper the level, the more realistic, right?!

Wrong! Including **more details** from lower levels leads to **more model parameters** which makes it harder to obtain a realistic model since the **realism** of a model is related to how **well-constrained** it is by experimental data.

- More model parameters means that more data is required to determine them, data which can often contain uncertainties and be hard to acquire.

Example: Describing the propagation of sound through air.

Large-Scale Modeling

Abstraction


- Modeler's wrench, **abstraction**: taking away aspects not important for answering the scientific questions which the model is designed to address, useful models can be formulated at different levels of organization without loss of tractability.
- Numerical simulations of brain network models have typically been based on either **abstract connectionist-type units** or **integrate-and-fire units**.
- **Subsampling** is often employed to decrease model size. With fewer pre-synaptic units providing synapses, it becomes necessary to exaggerate **connection density** or **synaptic conductance**. This results in a network with unnaturally few and strong signals circulating, in contrast to the real network, where many weak signals interact. Therefore differences might arise such as **artificial synchronization** which is a problem especially since synchronization is one of the more important phenomena.

Why Large-Scale Neuronal Networks?

- Improve understanding of brain functionality involving interactions of billions of neuronal and synaptic processes.
- Perform experiments (on a computer) that are impossible (experimentally or ethically) to be done on humans or animals.
- Eventually improve and test hypotheses about complex behaviors:
Perception • Attention • Learning • Memory • Consciousness • Sleep and wakefulness

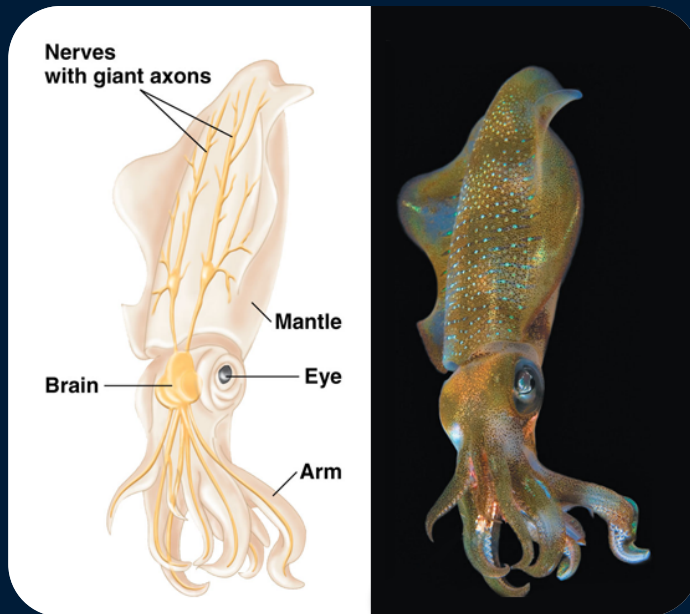
Large models need simple neuron models:

- Integrate-and-Fire types of models are obligatory because of their efficiency
- Izhikevich model is a wise choice because it exhibits a wide range of spiking behaviors and allows about 100 times faster computation runs than Hodgkin-Huxley





Neuron Models

Hodgkin Huxley



- Proposed model in 1952
- Based on voltage-clamp experiments on the squid giant axon
- Ion channels modeled as resistors and capacitors
- Membrane modeled as capacitor
- Received the 1963 Nobel Prize



$$C_m \frac{dV}{dt} = \underbrace{-g_L(V - V_L)}_{I_L} - \underbrace{\bar{g}_{Na} m^3 h (V - V_{Na})}_{I_{Na+}} - \underbrace{\bar{g}_K n^4 (V - V_K)}_{I_{K+}}$$

$$\frac{dm}{dt} = \alpha_m(V)(1 - m) - \beta_m(V)m$$


$$\frac{dh}{dt} = \alpha_h(V)(1 - h) - \beta_h(V)h$$

$$\frac{dn}{dt} = \alpha_n(V)(1 - n) - \beta_n(V)n$$

Three major currents:

- voltage-gated persistent K⁺ current I_{K+} with 4 activation gates (n_4)
- voltage-gated transient Na⁺ current I_{Na+} with 3 activation gates (m^3) and 1 inactivation gate (h)
- Ohmic leak current I_L (carried mostly by Cl⁻)

- Computationally complex
- Defined by 4 differential equations:
 - n controls potassium channel opening
 - m controls sodium channels opening
 - h controls sodium channels closing



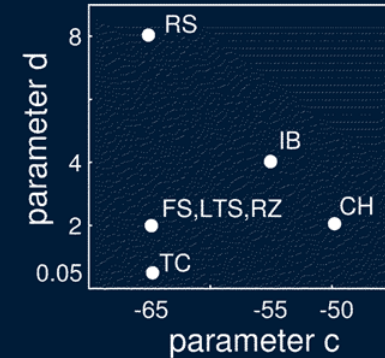
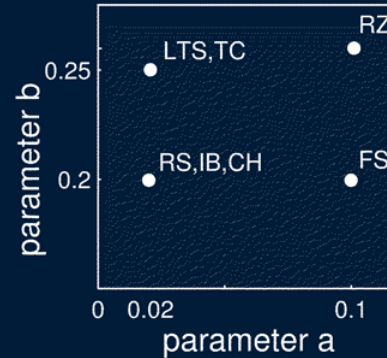
Neuron Models

Izhikevich's

- v represents the membrane potential
- u represents the membrane recovery variable (a negative feedback to v)
- I is a variable that represents synaptic currents

Parameters:

- a describes the time scale of the recovery variable u . Smaller values result in slower recovery (typically $a=0.02$)
- b describes the sensitivity of the recovery variable u to the sub-threshold fluctuations of the membrane potential v (typically $b=0.2$)
- c describes the after-spike reset value of the membrane potential v (typically $c=-65\text{mV}$)
- d describes the after-spike reset of the recovery variable u (typically $d=2$)




$$\frac{dv}{dt} = 0.04v^2 + 5v + 140 - u + I$$

$$\frac{du}{dt} = a(bv - u)$$

if $v = 30$ then $v = c$

$$u = u + d$$



Neuron Models

Izhikevich's

Implementation in Python

```
tmax = 1000.
dt = .5

neuron_model = 'RS'
if neuron_model=='RS':
    a = .02; b = .2; c = -65; d = 8.
elif neuron_model=='IB':
    a = .02; b = .2; c = -55; d = 4.
elif neuron_model=='FS':
    a = .1; b = .2; c = -65; d = 2.


Iapp = 7
tr = array([200., 700.])/dt

T = ceil(tmax/dt)
v = zeros(T)
u = zeros(T)

v[0] = -70
u[0] = -14

for t in arange(T-1):
    if t>tr[0] and t<tr[1]:
        I = Iapp
    else:
        I = 0
```

```
if v[t]<35:
    dv = (.04*v[t]+5)*v[t]+140-u[t]
    v[t+1] = v[t]+(dv+I)*dt
    du = a*(b*v[t]-u[t])
    u[t+1] = u[t]+dt*du
else:
    v[t] = 35
    v[t+1] = c
    u[t+1] = u[t]+d
```





- Spike Timing Dependent Plasticity (STDP)** is a temporally asymmetric form of Hebbian learning induced by tight temporal correlations between the spikes of pre- and post-synaptic neurons.
- Repeated pre-synaptic spike arrival a few milliseconds before postsynaptic action potentials leads in many synapse types to **long-term potentiation (LTP)** of the synapses, whereas repeated spike arrival after postsynaptic spikes leads to **long-term depression (LTD)** of the same synapse.
- The weight change Δw_j of a synapse from a presynaptic neuron j depends on the relative timing between presynaptic spike arrivals and postsynaptic spikes.
- t_j^f and t_i^f count pre-synaptic and post-synaptic spikes respectively:

$$\Delta w_j = \sum_{f=1}^N \sum_{n=1}^N W(t_i^n - t_j^f)$$

$$W(x) = A_+ \exp(-x/\tau_+) \quad \text{for } x > 0$$

$$W(x) = -A_- \exp(x/\tau_-) \quad \text{for } x < 0$$



Large-Scale Model

Sequential Code

Implementation in Python

```
from numpy import *
from pylab import *

Ne = 800 #excitatory neurons
Ni = 200 #inhibitory neurons
N = Ne+Ni #total neurons
M = 100 #synapses per neuron
D = 20 #maximal conduction delay
sm = 10.0 #maximal synaptic strength

class Izh(object):
    def __init__(self, a, d, v, u, init_s):
        self.a = a
        self.d = d
        self.I = 0
        self.v = v
        self.u = u
        self.post = None
        self.pre = []
        #synaptic weights
        self.s = [init_s] * M
        #derivatives of synaptic weights
        self.sd = [0.0] * M
        #distribution of delays
        self.delays = [[] for i in xrange(D)]
        #STDP functions
        self.LTP = [0.0]*(1001+D)
        self.LTD = 0.0
```

Disclaimer: This code is horribly slow! Though it is much more readable than the original C code, hence the implementation in Python for this presentation.

Large-Scale Model

Sequential Code

Implementation in Python (cont.)

```
#population of RS and FS type neurons
neurons = [Izh(.2, 8, -65, -13, 6.0) for i in xrange(Ne)] + \
    [Izh(.1, 2, -65, -13, -5.0) for i in xrange(Ni)]

#wiring neurons to each other
for i, neuron in enumerate(neurons):
    if i < Ne:
        neuron.post = list(permutation(N))[:M]
        for j in xrange(M):
            neuron.delays[randint(D)].append(j)
    else:
        neuron.post = list(permutation(Ne))[:M]
        neuron.delays[0] = range(M)

#find presynaptic neurons to each neuron
for neuron in neurons[:Ne]:
    for d in xrange(D):
        for syn in neuron.delays[d]:
            neurons[neuron.post[syn]].pre.append({'neuron':neuron, 'd':d, 'syn':syn})

firings = [[-D, 0]]
N_firings = 1
```

Large-Scale Model

Sequential Code

Implementation in Python (cont.)

```
for sec in xrange(60*60*24): #1 day
    for t in xrange(1000): #1 sec
        for neuron in neurons:
            neuron.I = 0 #reset input
        neurons[randint(N)].I = 20 #random thalamic input
        for i, neuron in enumerate(neurons):
            if neuron.v >= 30: #did it fire?
                neuron.v = -65.0
                neuron.u += neuron.d
                neuron.LTP[t+D] = .1
                neuron.LTD = .12
                for pre in neuron.pre:
                    #this spike was after pre-synaptic spikes
                    pre['neuron'].sd[pre['syn']] += pre['neuron'].LTP[t+D-pre['d']-1]
                firings.append([t, i])
                N_firings += 1
            for firing in reversed(firings[:N_firings-1]):
                if t-firing[0] >= D:
                    break
                for syn in neurons[firing[1]].delays[t-firing[0]]:
                    i = neurons[firing[1]].post[syn]
                    neurons[i].I += neurons[firing[1]].s[syn]
                    if firing[1]<Ne:
                        #this spike is before postsynaptic spikes
                        neurons[firing[1]].sd[syn] -= neurons[i].LTD
        for neuron in neurons:
            neuron.v += ((.04*neuron.v+5)*neuron.v+140-neuron.u+neuron.I)
            neuron.u += neuron.a*(.2*neuron.v-neuron.u)
            neuron.LTP[t+D+1] = .95*neuron.LTP[t+d]
            neuron.LTD *= .95
```


Large-Scale Model

Sequential Code

Implementation in Python (cont.)

```
print 'sec:', sec, 'firing rate:', float(N_firings)/N
#prepare for the next sec
for neuron in neurons:
    neuron.LTP[:D+1] = neuron.LTP[1000:1000+D+1]
k=N_firings-1
while 1000-firings[k][0]<D:
    k -= 1
for i in xrange(1, N_firings-k):
    firings[i][0] = firings[k+i][0]-1000
    firings[i][1] = firings[k+i][1]
N_firings = N_firings-k
#modify only exc connections
for neuron in neurons[:Ne]:
    for j in xrange(M):
        neuron.s[j] += .01+neuron.sd[j]
        neuron.sd[j] *= .9;
        if neuron.s[j]>sm:
            neuron.s[j] = sm
        elif neuron.s[j]<0:
            neuron.s[j] = 0.0
```

GPU Based Simulation

- The **GPU-SNN** model (running on an NVIDIA GTX-280 with 1GB of memory), is **up to 26 times faster** than a CPU version for the simulation of **100K neurons** with **50 Million synaptic connections**, firing at an average rate of 7Hz.
- For simulation of 100K neurons with 10 Million synaptic connections, the GPU-SNN model is only 1.5 times slower than real-time.

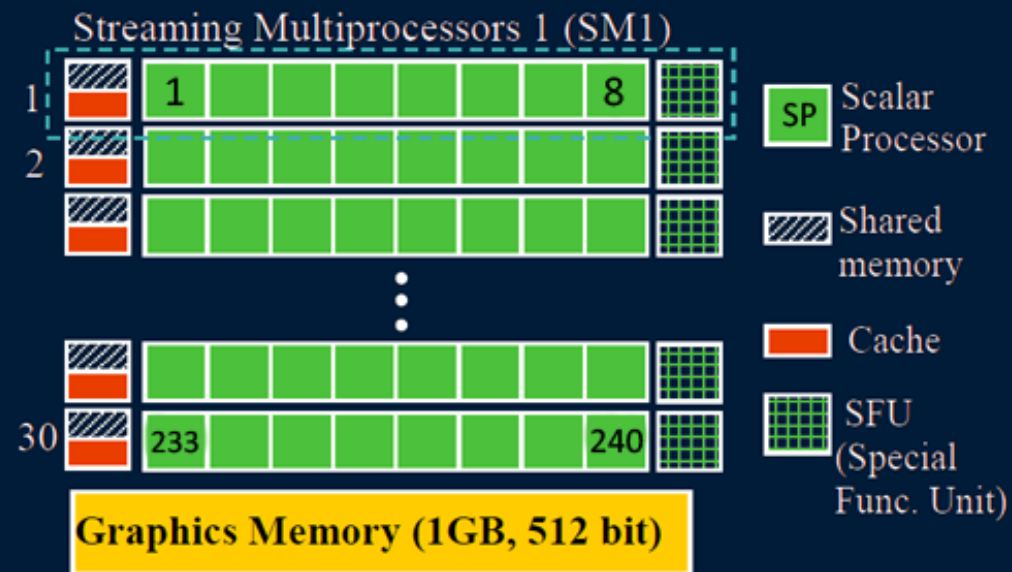
Challenges:

- Effective parallelism to **optimize the GPU resources** (processors, shared memory and memory bandwidth)
- Effective **handling of large fan-in/fan-out connections** to neurons
- Efficient usage of **limited GPU memory** for simulating large networks using sparse representations.

GPU Architecture

- Each **Streaming Multiprocessors (SM)** consists of eight floating-point Scalar Processors (SPs), a Special Function Unit (SFU), a multi-threaded instruction unit, a 16KB user-managed shared memory, and 16KB of cache memory.
- A single NVIDIA GTX280 GPU card consisting 240 scalar processors grouped into 30 SMs, each operating at 1.2 GHz, is used (350 GFLOPS).

- Each SM has a hardware thread scheduler that selects a group of threads, a.k.a **warp**, for execution.
- If any one of the threads in the group issues a costly external memory operation, then the thread scheduler automatically switches to a new thread group.



GPU Mapping

Parallelism Analysis

- **Neuronal parallelism (N-parallel)**: Each neuron is mapped on a processing element and computed in parallel. The synaptic computation for each neuron is carried out sequentially on its processing element. This mapping leads to warp divergence and is ineffective for GPUs.
- **Synaptic Parallelism (S-parallel)**: For a given neuron each synaptic connection is updated in parallel by different processing element. Thus synaptic information is distributed over all processing elements. The neuron computation is carried out sequentially. The maximum parallelism is limited by the number of synaptic connection that need to be updated in a given time step.
- **Neuronal-Synaptic Parallelism (NS-parallel)**: Uses both N-parallel and S-parallel techniques but at different stages in the simulation. At each time step where the neuron information needs to be updated, the N-parallel strategy is adopted. Thus, every thread within the GPU updates different neuron information in parallel. Whenever a spike is generated, the S-parallel mapping is deployed where the synapses need to be updated.

Large computation within diverging loop

```
i = threadIdx.x + blockIdx.x*blockDim.x
if ( membraneV[i] >= 30.0 ) {
    do_firing (i)           // 100-200 cycles
}                          // repeat for other neurons
```

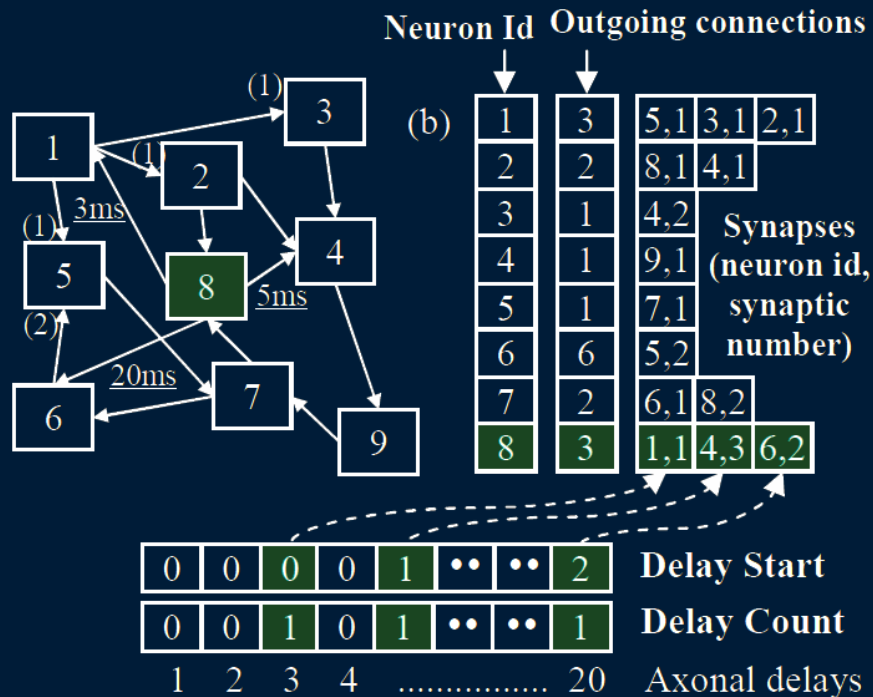
Warp divergence can occur if different threads within the same warp take different paths after a **branch condition**. If the diverging condition takes a large number of cycles, then other threads in the warp go into a busy waiting mode.

Small computation within diverging loop

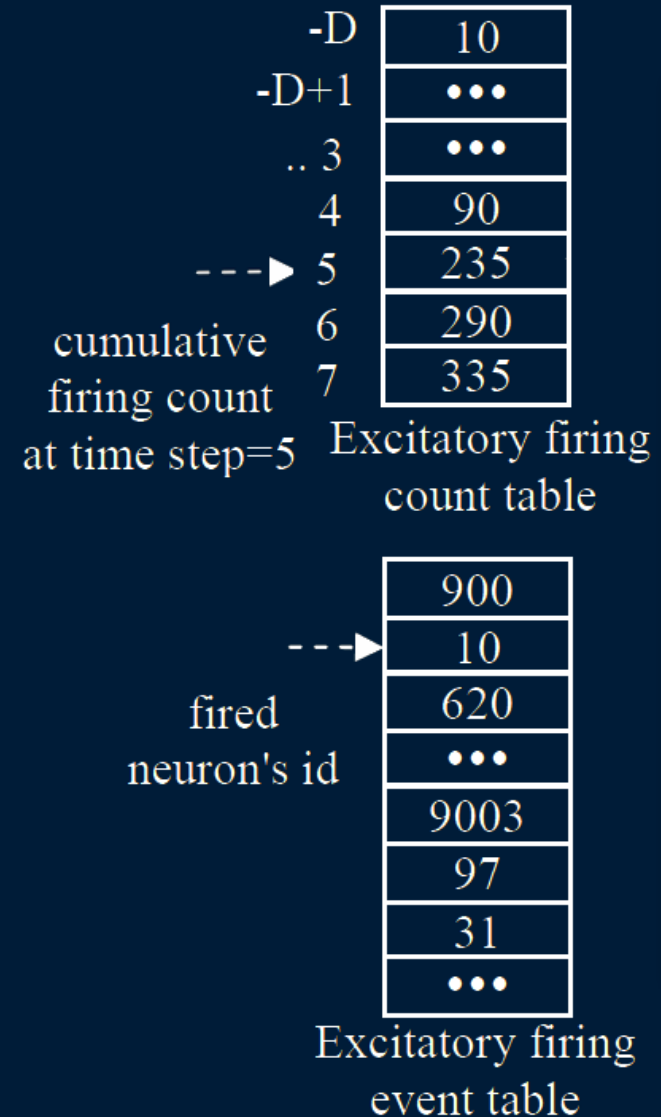
```
k=0; i = threadIdx.x + blockIdx.x*blockDim.x
if ( membraneV[i] >= 30.0 ) {
    p=atomicAdd(&k,1);buffer[p]=i //5-10 cycles
} // repeat for other neurons
__syncthreads();
offset = threadIdx.x;
while (offset < k)
    do_firing (buffer[offset])
    offset=offset+blockDim.x
```

GPU Based Simulation

Data Structure



Normally required memory is $O(NMD)$
 but we bring that down to $O(N(M+D))$
 with sparse representation.



GPU Based Simulation

Large Fan-in

Inputs: nid => Neuron Id,
I_fire => Input Fired Vector
s[i][j] => weight of i^{th} neuron, j^{th} synapses
len[i] => number of synaptic connections

Output: I_sum => Total synaptic current

Require: find_one[x] => pre-computed 256 entry table that returns the position of the first set bit in a given byte (e.g. find_one[0x10]=4, find_one[0x77]=0)

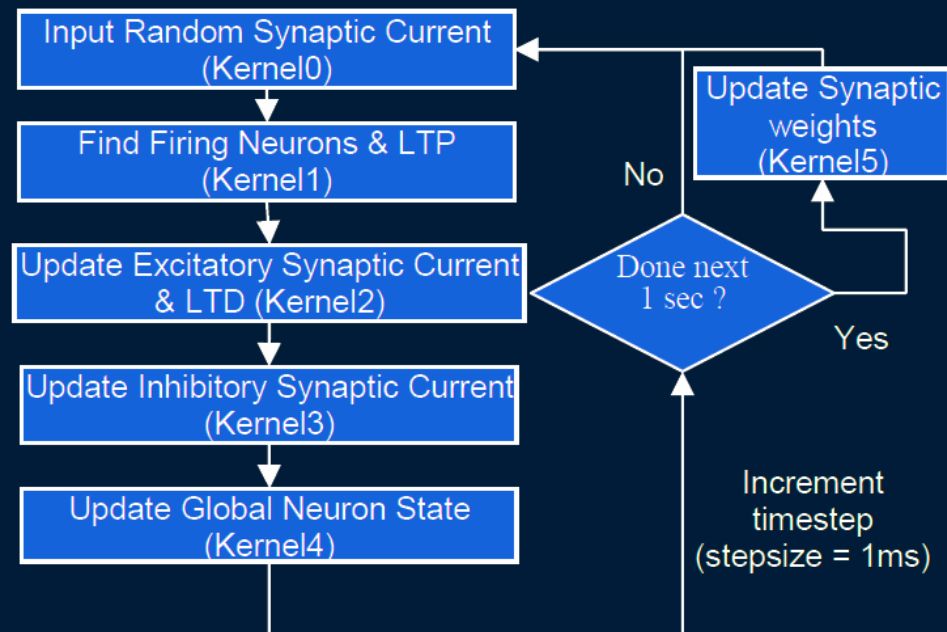
```
0. I_sum=0, y_end = ceil(len[i]/32)
1. for y=0:(y_end-1)
2.     part_I = read32(I_fire, y)    // Read  $y^{\text{th}}$  32 bit
                                     // from I_fire vector
3.     x = 0;
4.     while part_I ≠ 0
5.         byte_I = byte(part_I, x)  // Read  $x^{\text{th}}$  byte
6.         while byte_I ≠ 0
7.             idx = find_one[byte_I]
8.             set byte_I(idx) ← '0'
9.             I_sum = I_sum + s[nid][y*32+x*8+idx]
10.        part_I(x) ← 0; x = x+1;
11. return I_sum
```

- Large fan-ins of each neuron need to be calculated concurrently. Updating the synaptic current of post-synaptic neurons **atomically** is **infeasible** since due lack of atomic floating point operations in GPUs.
- Bit vector **I_fire** represents the input firing status of each neuron, whose up to 2-3 bits are set. The algorithm first scans the I_fire at the **word** level, then at the **byte** level and finally at **bit** level.
- If no bit is set, this approach incurs a small overhead of about 8 instruction cycles. This approach is memory efficient and has low computation overhead for a moderate number of input connections.

GPU Based Simulation

Simulation Flow


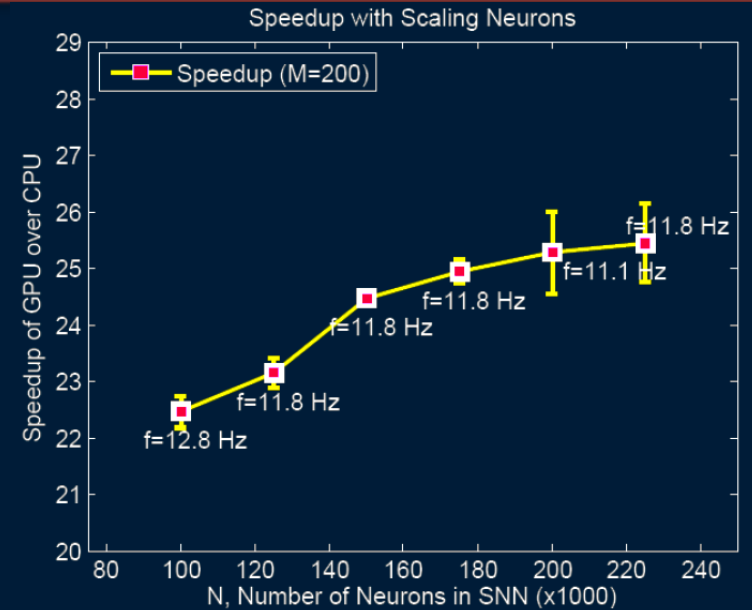
- For each kernel, various **block sizes** are tested in range of 30 to 120, with 128 threads in each block. The change in performance was very small for block sizes greater than 60.
- The network consist of N randomly connected excitatory (80%) and inhibitory (20%) neurons. For our experiments N ranges **from 50K to 225K neurons**. Number of synapses per neuron, M, ranges from 100 to 1000. The amount of change in the synaptic weight (using STDP rule) is accumulated during each time step; and **the weight is updated once a second** by Kernel5 (Figure 7) such that synaptic weight changes at a slower-rate than the neurons.



Results

- The speedup curves were obtained by **dividing the time taken by the CPU only mode and GPU mode** for simulating 10 seconds of model time (10,000 times steps with 1ms resolution) from the steady condition.
- We can observe that the overall speedup does not vary significantly for various values of N ($N > 105$). The variation in the speedup curve is mainly **due to the variation in the firing rate**. An increase in the firing rate causes slight improvement in the speedup.
- For $M=100$ and $N=105$, the speedup is limited to 18. The GPU takes **15 seconds to simulate 10 seconds** of model time. For larger values of M the speedup jumps from 18 to around 25 **due to increases in the available synaptic parallelism**.

Scalability Analysis



Results

Fidelity Test

- The **GPU model differs from the reference model** in the following ways: implementation of STDP calculations, network representation, firing information representation, etc.
- Thus **direct comparison is difficult** because the SNN state can change significantly even if one spike is altered.
- To ensure the accuracy and fidelity of GPU implementation various neuronal metrics are considered:
 - difference in average firing rate
 - difference in the synaptic weights of excitatory connections
 - difference in the inter-spike intervals (ISI) for excitatory neurons and for inhibitory neurons

Metrics	N=1000, M=100	N=3000, M=100
Synaptic Weights	0.992	0.099
ISI (Excitatory)	0.799	0.144
ISI (Inhibitory)	0.677	0.261


Comparison of distribution of synaptic weights and ISI

Firing Rate Metrics (Hz)	N=1000, M=100		N=3000, M=100	
	Matlab	GPU	Matlab	GPU
Excitatory Neurons	3.1423 (0.4934)	3.1693 (0.7034)	3.8242 (0.1413)	3.3855 (0.0843)
Inhibitory Neurons	24.95 (3.3683)	22.0345 (4.5293)	31.5863 (1.0140)	24.9593 (0.5713)

Comparison of firing rate between MATLAB and GPU implementations

Thanks for your attention :)

Any Questions?!



Contact me at nima.mohammadi@ut.ac.